

# Unix System Programming:

## The Big Picture

류은경  
ekryu@knu.ac.kr

# Unix System Programming

---

## **1.2 What Is System Programming?**

1.3 Understanding System Programming

1.4 Unix from the User Perspective

1.5 Unix from the Larger Perspective

1.6 Can I Try to Do One?

# The Simple Program Model

- ♦ There are many sorts of programs; many programs are based on the model below:




FIGURE 1.1  
An application program in a computer.

## ◆ Typical Program in this Model

```
/* copy from stdin to stdout */  
main()  
{  
    int c;  
    while( ( c = getchar() ) != EOF )  
        putchar(c);  
}
```




FIGURE 1.2

How application programs see user I/O.

# Reality

- ♦ Wha if you log into a **multiuser system**?
  - like a typical Unix machine





FIGURE 1.3

Reality: many users, programs, and devices.



**FIGURE 1.4**

How are they all connected?

# Operating System

- ♦ To manage and protect all the resources
- ♦ To **connect** the various *devices* to the *programs*




FIGURE 1.5

An operating system is a program.

# Providing Services to Programs

- ♦ **The kernel is the only program with access to devices**
  - Programs request *services* from the OS
  - The OS contains *code* to provide services




FIGURE 1.6

The kernel manages all connections.



# Unix System Programming

---

1.2 What Is System Programming?

## **1.3 Understanding System Programming**

1.4 Unix from the User Perspective

1.5 Unix from the Larger Perspective

1.6 Can I Try to Do One?

## ♦ The services provided by the kernel :

- Processors: ...
- Input / Output: ...
- Process management: ...
- Memory: ...
- Devices: ...
- Timers: ...
- Inter-process communication: ...
- Networking: ...

# Understanding Systems Programming

---

- ◆ **Understanding how the kernel works and how to write programs that use these services**
  - What are the details of each type of the kernel service?
  - How does data get from a device to the program and back?
- ◆ **System programs use those services directly**
  - Also known as **system utilities**
  - Provide a convenient environment for program development and execution

# Our Method: Three Simple Steps

- ◆ We shall learn about *Unix services* by

1. Looking at “real” programs : `ls`, `cat`, `grep` ...
2. Looking at the **system calls** to invoke the services : `fork()`, `exec()`, `open()` ...
3. Writing our own version




Figure 1.1 Architecture of the UNIX operating system

# Unix System Programming

---

1.2 What Is System Programming?

1.3 Understanding System Programming

**1.4 Unix from the User Perspective**

1.5 Unix from the Larger Perspective

1.6 Can I Try to Do One?

# What does Unix do?


---

- ◆ **Log In – Run Programs – Log Out**
- ◆ **Working with Directories**
  - mkdir, rmdir, cd, ls
- ◆ **Working with Files**
  - touch, rm, cp, mv

# Log In – Run Programs – Log out

## ♦ How Does That work?

- Login: ...
- Run programs: .....
- Logout: ...



```
Linux 1.2.13 (maya) (ttyp1)
```

```
maya login: betsy
```

```
Password: _
```

```
$
```

```
$ date
```

```
Sat Jul 1 21:34:10 EDT 2000
```

```
$ _
```

```
$ exit
```

FIGURE 1.7

A user logged into a computer.

# Working with Directories (1/2)

- ♦ Unix organizes *files* into a tree-structured dir system
- ♦ A Tree of Directories :

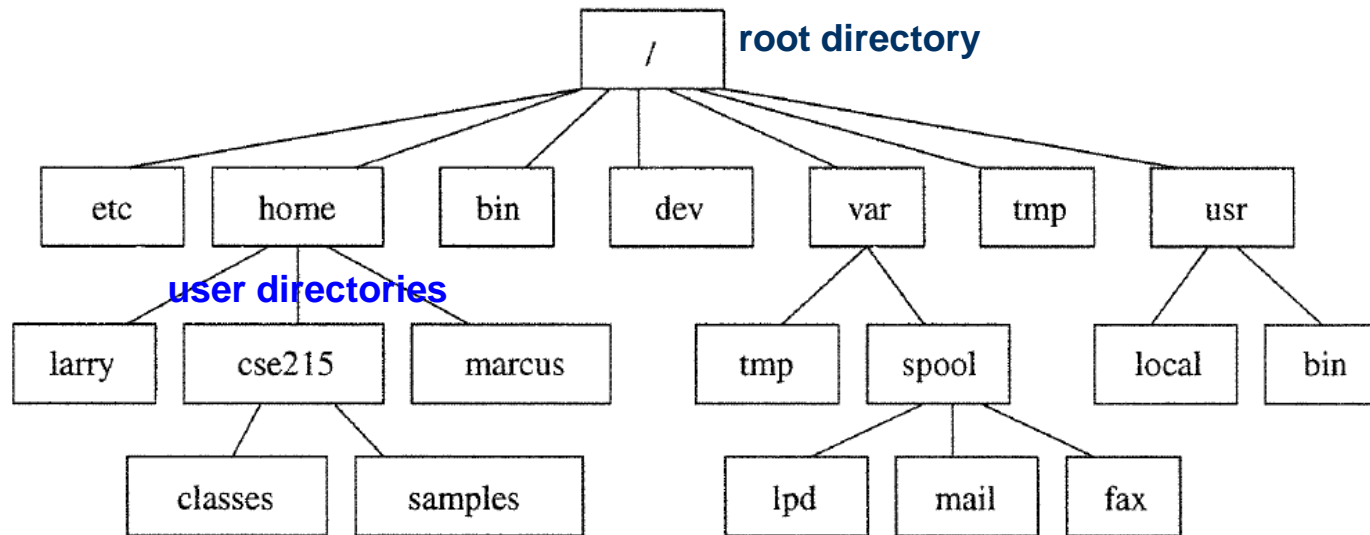


FIGURE 1.8

Part of the directory tree.



# Working with Directories (2/2)

## ◆ Commands for Working with Directories

- **ls** – list directory contents

```
$ ls /etc
$ ls /
```
- **cd** – change to a different directory

```
$ cd /bin
$ cd ..
$ cd
```
- **pwd** – print path to current directory

```
$ pwd
```
- **mkdir, rmdir** – make and remove directories

```
$ cd
$ mkdir jokes
$ rmdir jokes
```

## Working with Files (1/2)

- ◆ **Names of Files:** up to about 255 characters
- ◆ **cat, more, less – examine file contents**

```
$ cat shopping-list
soap
cornflakes
milk
apples
jam
$
$ more longfile
```

- ◆ **cp – make a copy of a file**

```
$ cp shopping-list last.week.list
```

## Working with Files (2/2)

- ♦ **rm – delete a file**

```
$ rm old.data junk shopping.june1992
```

- ♦ **mv – rename or move a file**

```
$ mv prog1.c first_program.c
```

Rename

```
$ mkdir mycode
```

```
$ mv first_program.c mycode
```

Move

- ♦ **lpr, lp – print file on paper**

```
$ lpr filename
```

# File Permission Attributes

- ◆ **To allow users to control access to their files :**

- Unix assigns to each file several attributes

- ◆ **ls -l** shows attributes of a file:

```
$ ls -l outline.01
```

```
-rwxr-x---    1 molay    users    1064 Jun 29 00:39 outline.01
```

file	user	group	size	modification date and
permission	name	name		time
attributes	(owner/ creator)			



- r w x	r w x	r w x	r: read, w: write, x:execute
user	group	other	

# Unix System Programming

---

1.2 What Is System Programming?

1.3 Understanding System Programming

1.4 Unix from the User Perspective

**1.5 Unix from the Larger Perspective**

1.6 Can I Try to Do One?

# Larger Perspective

---

- ♦ The *larger system* is one that consists of
  - more than one user,
  - more than one program,
  - more than one computer, and
  - the connections among these people, programs, and computers.
- ♦ Ex)
  - Internet bridge tournaments
  - bc/dc program

# Internet Bridge Tournaments

- ◆ This bridge example introduces the *three topics* in Unix system programming:

- Communication: ...
- Process Coordination: ...
- Network Access: ..




FIGURE 1.11

Separate programs send messages to each other.



Online Bridge

# bc Calculator

```
$ bc
```

```
2+3*4+5*10
```

```
64
```

```
999*888
```

```
887112
```

```
333^44
```

```
97173148450073933242567011681978355452978308136572718207376696208658\
```

```
8054474264529237756850460517373752303625521
```

```
x=3
```

```
if (x==3) {
```

```
y=x*3;
```

```
}
```

```
y
```

```
9
```

```
$
```

To exit from bc, press **Ctrl-D** or quit



```
echo Hello, World!  
echo "4 * (2 + 3)" | bc  
echo "This is a test" > test.txt
```

```
$ bc
2 + 3
5
<-- press Ctrl-Z here
```

Stopped

\$ **ps**      ※ lists the processes you are running

PID	TTY	S	TIME	CMD
25102	ttyp2	T	0:00.02	bc
27081	ttyp2	T	0:00.01	dc -
27560	ttyp2	I	0:00.59	-bash
27681	ttyp2	T	0:00.00	bc

```
$ fg
<-- press Ctrl-D here
```

The **fg** continues a stopped job by running it in the foreground.

```
$ man dc
```

```
User Commands dc(1)
```

#### NAME

```
dc - desk calculator
```

#### SYNOPSIS

```
dc [ filename ]
```

#### DESCRIPTION

dc is an arbitrary precision arithmetic package. Ordinarily it operates on decimal integers, but one may specify an input base, output base, and a number of fractional digits to be maintained. The overall structure of dc is a stacking (reverse Polish) calculator. If an argument is given, input is taken from that file until its end, then from the standard input.

**\$dc** ※ postfix notation

2

3

+

p

5

# bc Calculator

- ♦ It communicates with a process running `dc`
  - through a communication system called *pipes*

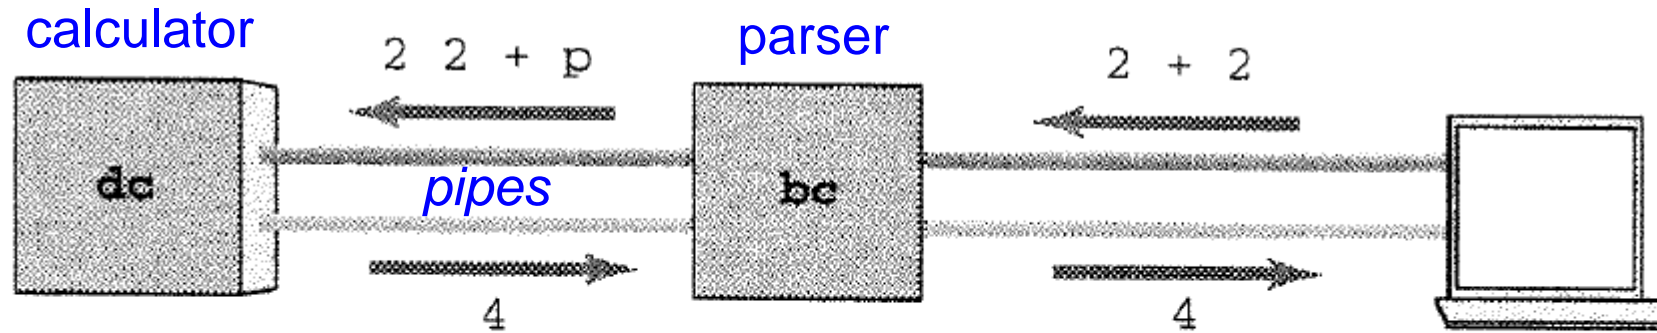


FIGURE 1.12

Programs send messages to each other.

※ The GNU version of `bc` uses an internal stack-based calculator instead of `dc`.

- ◆ **This bc/dc is another program that involves**
  - different processes
  - some sort of communication
  - cooperation
- ◆ **Learning system programming consists of learning**
  - how to build these separate programs and
  - how to build the connections and cooperation.

# Unix System Programming

---

1.2 What Is System Programming?

1.3 Understanding System Programming

1.4 Unix from the User Perspective

1.5 Unix from the Larger Perspective

**1.6 Can I Try to Do One?**

# “What does more do?” “How does more work?”

- ◆ It displays a file one screenful at a time

```
$ more /usr/include/utmp.h
```

- ◆ USAGE:

```
$ more filename  
$ command | more  
$ more < filename
```

- ◆ Logic

```
+----> show 24 lines from input  
| +--> print [more?] message  
| |   Input Enter, SPACE, or q  
| +--  if Enter, advance one line  
+----  if SPACE  
      if q --> exit
```

## 1<sup>st</sup> version : more01.c

---

### ◆ Compile and run it

```
$ cc more01.c -o more01
$ more01 more01.c
```

< Linux >  
※ cc → gcc  
※ \$ ./more01 more01.c



```
/* more01.c - version 0.1 of more
 *      read and print 24 lines then pause for a few special commands
 */
```

```
#include <stdio.h>
#include <stdlib.h>
```

```
#define PAGELEN 24
```

```
#define LINELEN 512
```

```
void do_more(FILE *);
```

```
int see_more();
```

```
int main( int ac , char *av[] )
```

```
{
    FILE *fp;
    if ( ac == 1 )
        do_more( stdin );
    else
        while ( --ac )
            if ( (fp = fopen( *++av , "r" )) != NULL )
            {
                do_more( fp ) ;
                fclose( fp );
            }
            else
                exit(1);

    return 0;
}
```

**\$ ./more01 filename**

```

void do_more( FILE *fp )
/*
 * read PAGELEN lines, then call see_more() for further instruction
 */
{
    char    line[LINELLEN];
    int     num_of_lines = 0;
    int     see_more(), reply;

    while ( fgets( line, LINELLEN, fp ) ){
        if ( num_of_lines == PAGELEN ) {
            reply = see_more();
            if ( reply == 0 )
                break;
            num_of_lines -= reply;
        }
        if ( fputs( line, stdout ) == EOF )
            exit(1);
        num_of_lines++;
    }
}

```

```

/* more input more?
/* full screen? */
/* y: ask user */
/*    n: done    */

/* reset count */

/* show line */
/* or die */
/* count it */

```

```

/* more01.c - version 0.1 of more
 *          read and print 24 lines then pause for a few special commands
 */

#include <stdio.h>
#include <stdlib.h>

#define PAGELEN 24
#define LINELLEN 512

void do_more(FILE *);
int see_more();

int main( int ac, char *av[] )
{
    FILE *fp;
    if( ac == 1)
        do_more( stdin );
    else
        while( --ac )
            if(( fp = fopen( *++av, "r" )) != NULL )
            {
                do_more( fp );
                fclose( fp );
            }
}

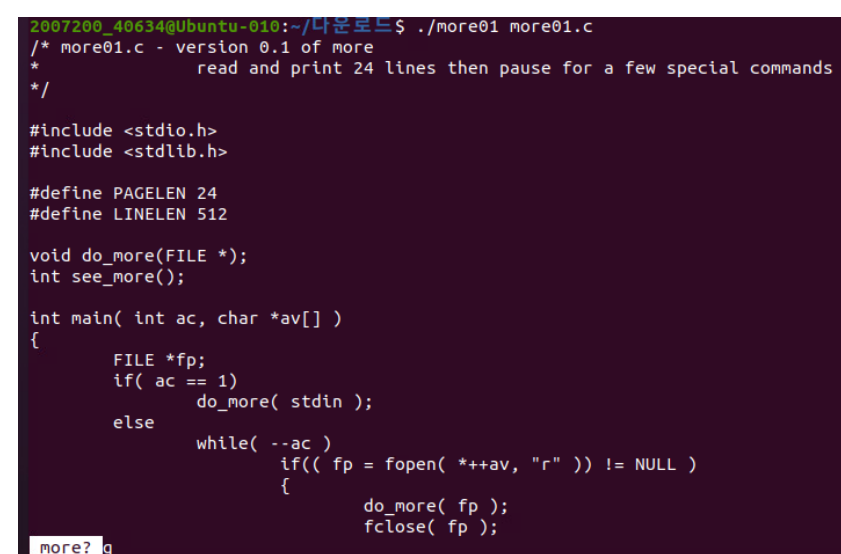
```

```

int see_more()
/*
 *   print message, wait for response, return # of lines to advance
 *   q means no, space means yes, CR means one line
 */
{
    int        c;

    printf("\033[7m more? \033[m");
    while( (c=getchar()) != EOF )
    {
        if ( c == 'q' )
            return 0;
        if ( c == ' ' )
            return PAGELEN;
        if ( c == '\n' )
            return 1;
    }
    return 0;
}

```



```

2007200_40634@Ubuntu-010:~/나온보드$ ./more01 more01.c
/* more01.c - version 0.1 of more
 *           read and print 24 lines then pause for a few special commands
 */

#include <stdio.h>
#include <stdlib.h>

#define PAGELEN 24
#define LINELEN 512

void do_more(FILE *);
int see_more();

int main( int ac, char *av[] )
{
    FILE *fp;
    if( ac == 1 )
        do_more( stdin );
    else
        while( --ac )
            if( ( fp = fopen( *++av, "r" ) ) != NULL )
            {
                do_more( fp );
                fclose( fp );
            }
}

```

```

+----> show 24 lines from input
| +--> print [more?] message
| |   Input Enter, SPACE, or q
| +-- if Enter, advance one line
+---- if SPACE
      if q --> exit

```

# 1<sup>st</sup> version : more

## ◆ Subtle problems:

- If you press the space bar or the “q” key, nothing happens **until you press Enter**.
- Also, the little “more?” message is still there.

```
2007200_40634@Ubuntu-010:~/다운로드$ ./more01 more01.c
/* more01.c - version 0.1 of more
 *          read and print 24 lines then pause for a few special commands
 */

#include <stdio.h>
#include <stdlib.h>

#define PAGELEN 24
#define LINELEN 512

void do_more(FILE *);
int see_more();

int main( int ac, char *av[] )
{
    FILE *fp;
    if( ac == 1 )
        do_more( stdin );
    else
        while( --ac )
            if(( fp = fopen( *++av, "r" )) != NULL )
            {
                do_more( fp );
                fclose( fp );
            }
}
```

more? q

## ◆ Using Pipeline

```
$ who | more
```

```
$ ls /bin | ./more01
```

```
$ ls /bin | more
```

```
$ ls /bin | ./more01
```




FIGURE 1.14

more reads stdin.

- **more01** does **not pause** after 24 lines;
- Reason? ...

◆ How does the real `more` solve this problem?

- Read from the keyboard directly!
- There is a special file, called `/dev/tty`, which is actually a connection to the keyboard and screen.

```
$ ls /bin | more
```

```
$ ls /bin | ./more01
```




FIGURE 1.15

`who` reads user input from a terminal.

```
/* more02.c - version 0.2 of more
 *      read and print 24 lines then pause for a few special commands
 *      feature of version 0.2: reads from /dev/tty for commands
 */
```

```
#include      <stdio.h>
#include      <stdlib.h>
```

```
#define PAGELEN 24
```

```
#define LINELEN 512
```

```
void do_more(FILE *);
```

```
int see_more(FILE *);
```

```
int main( int ac , char *av[] )
```

```
{
```

```
    FILE *fp;
```

```
    if ( ac == 1 )
```

```
        do_more( stdin );
```

```
    else
```

```
        while ( --ac )
```

```
            if ( (fp = fopen( *++av , "r" )) != NULL )
```

```
            {
```

```
                do_more( fp ) ;
```

```
                fclose( fp );
```

```
            }
```

```
        else
```

```
            exit(1);
```

```
    return 0;
```

```
}
```

```

void do_more( FILE *fp )
/*
 * read PAGELEN lines, then call see_more() for further instructions
 */
{
    char    line[LINELLEN];
    int     num_of_lines = 0;
    int     see_more(FILE *), reply;
    FILE    *fp_tty;

    fp_tty = fopen( "/dev/tty", "r" );           /* NEW: cmd stream */
    if ( fp_tty == NULL )                        /* if open fails */
        exit(1);                                /* no use in running*/

    while ( fgets( line, LINELLEN, fp ) ){       /* more input */
        if ( num_of_lines == PAGELEN ) {         /* full screen? */
            reply = see_more(fp_tty);            /* NEW: pass FILE */
            if ( reply == 0 )                     /* n: done */
                break;
            num_of_lines -= reply;                /* reset count */
        }
        if ( fputs( line, stdout ) == EOF )      /* show line */
            exit(1);                             /* or die */
        num_of_lines++;                          /* count it */
    }
}

```



```

int see_more(FILE *cmd)                                /* NEW: accepts arg */
/*
 * print message, wait for response, return # of lines to advance
 * q means no, space means yes, CR means one line
 */
{
    int    c;
    printf("\033[7m more? \033[m");                    /* reverse on a vt100 */
    while( (c=getc(cmd)) != EOF )                      /* NEW: reads from tty */
    {
        if ( c == 'q' )                                /* q -> N */
            return 0;
        if ( c == ' ' )                                /* ' ' => next page */
            return PAGELEN;                             /* how many to show */
        if ( c == '\n' )                                /* Enter key => 1 line */
            return 1;
    }
    return 0;
}

```

## 2<sup>nd</sup> version : more

### ◆ Compile and run it

```
$ cc -o more02 more02.c
```

```
$ ls /bin | ./more02
```




FIGURE 1.16

The connection to the terminal has settings.

```
2007200_40634@Ubuntu-010:~/다운로드$ ls /bin | ./more02
2to3-2.7
GET
HEAD
POST
VGAAuthService
X
X11
Xephyr
Xorg
Xwayland
[
aa-enabled
aa-exec
aconnect
acpi_listen
add-apt-repository
addpart
addr2line
alsabat
alsaloop
alsamixer
alsatplg
alsaucm
amidi
more?
```




FIGURE 1.17

A diagram of the main structure of a Unix system.

# SUMMARY

---

- ♦ **Computer systems that run several programs for several users at the same time require a central management program.**
- ♦ **The Unix kernel is a program that schedules programs and controls access to resources.**
  - User programs ask the kernel for access to resources.
  - Some Unix programs consist of separate programs that share or exchange data.
- ♦ **Writing systems programs requires an understanding of the structure and use of kernel services.**