

Users, Files, and the Manual : who

류은경
ekryu@knu.ac.kr

Objectives

◆ Ideas and Skills

- The role and use of on-line documentation
- The Unix file interface: `open`, `read`, `write`, `lseek`, `close`
- Reading, creating, and writing files
- File descriptors
- Buffering: user level and kernel level
- Kernel mode, user mode, and the cost of system calls
- How Unix represents time, how to format Unix time
- Using the `utmp` file to find list of current users
- Detecting and reporting errors in system calls

◆ System Calls and Functions

- `open`, `read`, `write`, `create`, `lseek`, `close`
- `perror`

◆ Commands

- `man`, `who`, `cp`, `login`

2.2 Asking about who

2.3 What Does who Do?

2.4 How Does who Do It?

2.5 Can I Write who?

2.6 Writing cp (read and write)

2.7 More Efficient File I/O: Buffering

2.8 Buffering and the Kernel

2.10 What to Do with System-Call Errors

Commands are Programs

- ◆ Almost all Unix **commands** are simply **programs** written by a variety of people, usually in C.
- ◆ **Adding new commands** to Unix:
 - You write a new program and have the executable file stored in one of the standard system directories:
`/bin, /usr/bin, /usr/local/bin.`

```
ls  
cd  
pwd  
cp  
mv  
rm  
mkdir  
touch  
man  
echo  
...
```

Q1: What Does who Do?

```
$ who
heckerl      ttyp1      Jul 21 19:51    (tide75.surfcity.com)
nlopez       ttyp2      Jul 21 18:11    (roam163-141.student.ivy.edu)
dgsulliv     ttyp3      Jul 21 14:18    (h004005a8bd64.ne.mediaone.net)
ackerman     ttyp4      Jul 15 22:40    (asd1-254.fas.state.edu)
wwchen       ttyp5      Jul 21 19:57    (circle.square.edu)
barbier      ttyp6      Jul  8 13:08    (labpc18.elsie.special.edu)
ramakris     ttyp7      Jul 13 08:51    (roam157-97.student.ivy.edu)
czhu         ttyp8      Jul 21 12:47    (spa.sailboat.edu)
bpsteven     ttyp9      Jul 21 18:26    (207.178.203.99)
molay        ttypa      Jul 21 20:00    (xyz73-200.harvard.edu)
```

\$ ↑ ↑ ↑ ↑ one log-in session

username terminal name login time login host

Reading the Manual

```
$ man who
```

```
who(1)
```

```
NAME
```

```
who - Identifies users currently logged in
```

```
SYNOPSIS
```

```
who [-a] | [-AbdhHlmMpqrstTu] [file]
```

```
who am i
```

```
who am I
```

```
whoami
```

```
...
```

Q2: How Does who Do It?

◆ To learn more about Unix from Unix

- Read the manual
- Search the manual
- Read the .h files
- Follow SEE ALSO links

Read the Manual

\$man who

DESCRIPTION

The who utility can list the user's name, terminal line, login time, elapsed time since activity occurred on the line, and the process-ID of the command interpreter (shell) for each current UNIX system user. It examines the /var/adm/utmp file to obtain its information. If file is given, that file (which must be in utmp(4) format) is examined. Usually, file will be /var/adm/wtmp, which contains a history of all the logins since the file was last created.

Search the Manual

♦ \$ man -k utmp

endutent	getutent (3c)	- access utmp file entry
endutxent	getutxent (3c)	- access utmpx file entry
getutent	getutent (3c)	- access utmp file entry
getutid	getutent (3c)	- access utmp file entry
getutline	getutent (3c)	- access utmp file entry
getutmp	getutxent (3c)	- access utmpx file entry
getutmpx	getutxent (3c)	- access utmpx file entry
getutxent	getutxent (3c)	- access utmpx file entry
getutxid	getutxent (3c)	- access utmpx file entry
getutxline	getutxent (3c)	- access utmpx file entry
pututline	getutent (3c)	- access utmp file entry
pututxline	getutxent (3c)	- access utmpx file entry
setutent	getutent (3c)	- access utmp file entry
setutxent	getutxent (3c)	- access utmpx file entry
ttyslot	ttyslot (3c)	- find the slot in the utmp file of the current user
updwtmp	getutxent (3c)	- access utmpx file entry
updwtmpx	getutxent (3c)	- access utmpx file entry
<u>utmp</u>	<u>utmp (4)</u>	<u>- utmp and wtmp entry formats</u>
utmp2wtmp	acct (1m)	- overview of accounting and miscellaneous accounting commands
utmpd	utmpd (1m)	- utmp and utmpx monitoring daemon
utmpname	getutent (3c)	- access utmp file entry
utmpx	utmpx (4)	- utmpx and wtmpx entry formats
utmpxname	getutxent (3c)	- access utmpx file entry
<u>wtmp</u>	<u>utmp (4)</u>	<u>- utmp and wtmp entry formats</u>
wtmpx	utmpx (4)	- utmpx and wtmpx entry formats
\$		

\$ **man 4 utmp**

utmp(4)

utmp(4)

NAME

utmp, wtmp - Login records

SYNOPSIS

```
#include <utmp.h>
```

DESCRIPTION

The utmp file records information about who is currently using the system.

The file is a sequence of utmp entries, as defined in struct utmp in the utmp.h file.

The utmp structure gives the name of the special file associated with the user's terminal, the user's login name, and the time of the login in the form of time(3). The ut_type field is the type of entry, which can specify several symbolic constant values. The symbolic constants are defined in the utmp.h file.

The wtmp file records all logins and logouts. A null user name indicates a logout on the associated terminal. A terminal referenced with a tilde (~) indicates that the system was rebooted at the indicated time. The adjacent pair of entries with terminal names referenced by a vertical bar (|) or a right brace (}) indicate the system-maintained time just before and just after a date command has changed the system's time frame.

The wtmp file is maintained by login(1) and init(8). Neither of these pro-grams creates the file, so, if it is removed, record keeping is turned off. See ac(8) for information on the file.

FILES

/usr/include/utmp.h

/var/adm/utmp

more (88%)

Read the .h files

♦ \$ more /usr/include/utmp.h

```
#define UTMP_FILE      "/var/adm/utmp"
#define WTMP_FILE      "/var/adm/wtmp"

#include <sys/types.h> /* for pid_t, time_t */

/*
 * Structure of utmp and wtmp files.
 *
 * Assuming these numbers is unwise.
 */

#define ut_name ut_user /* compatibility */
struct utmp {
    char    ut_user[32]; /* User login name */
    char    ut_id[14]; /* /etc/inittab id- IDENT_LEN in
                       * init */
    char    ut_line[32]; /* device name (console, lnxx) */
    short   ut_type; /* type of entry */
    pid_t   ut_pid; /* process id */
    struct exit_status {
        short   e_termination; /* Process termination status */
        short   e_exit; /* Process exit status */
    } ut_exit; /* The exit status of a process
               * marked as DEAD_PROCESS.
               */
    time_t   ut_time; /* time entry was made */
    char     ut_host[64]; /* host name same as
                          * MAXHOSTNAMELEN */
};
/* Definitions for ut_type */

utmp.h (60%)
```

```
struct utmp {
    short   ut_type; /* Type of record */
    pid_t   ut_pid; /* PID of login process */
    char    ut_line[UT_LINESIZE]; /* Device name of tty - "/dev/" */
    char    ut_id[4]; /* Terminal name suffix,
                     * or inittab(5) ID */
    char    ut_user[UT_NAMESIZE]; /* Username */
    char    ut_host[UT_HOSTSIZE]; /* Hostname for remote login, or
                                   kernel version for run-level
                                   messages */
    struct exit_status ut_exit; /* Exit status of a process
                                * marked as DEAD_PROCESS; not
                                * used by Linux init (1 */
    /* The ut_session and ut_tv fields must be the same size when
       compiled 32- and 64-bit. This allows data files and shared
       memory to be shared between 32- and 64-bit applications. */
    #if __WORDSIZE == 64 && defined __WORDSIZE_COMPAT32
        int32_t ut_session; /* Session ID (getsid(2)),
                             * used for windowing */
        struct {
            int32_t tv_sec; /* Seconds */
            int32_t tv_usec; /* Microseconds */
        } ut_tv; /* Time entry was made */
    #else
        long   ut_session; /* Session ID */
        struct timeval ut_tv; /* Time entry was made */
    #endif
    int32_t ut_addr_v6[4]; /* Internet address of remote
                           * host; IPv4 address uses
                           * just ut_addr_v6[0] */
    char __unused[20]; /* Reserved for future use */
};
```

\$ man 5 utmp

\$ utmpdump /var/run/utmp

How who works




FIGURE 2.2


Data flow in the who command.

Can I Write who?

- ◆ ***Two tasks we need to program***
 - Read structs from a file
 - Display the information stored in a struct

Q: How do I read structs from a file?

◆ Let's Read the Manual!



```
$ man -k file | grep read
_llseek (2)          - reposition read/write file offset
fileevent (n)        - Execute a script when a channel becomes readable
                      or writable
gftype (1)           - translate a generic font file for humans to read
lseek (2)            - reposition read/write file offset
macsave (1)          - Save Mac files read from standard input
read (2)           - read from a file descriptor
readprofile (1)      - a tool to read kernel profiling information
scr_dump, scr_restore, scr_init, scr_set (3) - read (write) a curses
screen from (to) a file
tee (1)              - read from standard input and write to standard
                      output and files
$
```

\$ man 2 read

READ(2)

System calls

READ(2)

NAME

read - read from a file descriptor

SYNOPSIS

#include <unistd.h>

ssize_t read(int fd, void *buf, size_t count);

DESCRIPTION

read() attempts to read up to count bytes from file descriptor fd into the buffer starting at buf.

RELATED INFORMATION (called SEE ALSO in some versions)

Functions: fcntl(2), creat(2), dup(2), ioctl(2), getmsg(2), lockf(3), lseek(2), mtio(7), open(2), pipe(2), poll(2), socket(2), socketpair(2), termios(4), streamio(7), opendir(3) lockf(3)

Standards: standards(5)

ANS: we use `open`, `read`, and `close`

◆ Opening a file: `open()`

- It is a **kernel service**; system call

```
fd = open(name, mode)
```

 ↑ ↙
char * O_RDONLY
 O_WRONLY
 O_RDWR

returns -1 on error
OR an int on success

```
int fd = open("file.txt", O_RDONLY);  
if (fd == -1) {  
    perror("open");  
    return -1;  
}  
  
close(fd);
```




FIGURE 2.3

A file descriptor is a connection to a file.


```
1 #include <stdio.h>
2 #include <fcntl.h>
3 #include <unistd.h>
4
5 int main() {
6     int fd = open("file_open.c", O_RDONLY);
7     if (fd == -1) return 1;
8
9     char buffer[100];
10    ssize_t bytesRead = read(fd, buffer, sizeof(buffer) - 1);
11
12    if (bytesRead > 0) {
13        buffer[bytesRead] = '\0';
14        printf("%s\n", buffer);
15    }
16
17    close(fd);
18    return 0;
19 }
```

file_open.c

open		
PURPOSE	Creates a connection to a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = open(char *name, int how)	
ARGS	name	name of file
	how	O_RDONLY, O_WRONLY, or O_RDWR
RETURNS	-1	on error
	int	on success

read		
PURPOSE	Transfer up to qty bytes from fd to buf	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t numread = read(int fd, void *buf, size_t qty)	
ARGS	fd	source of data
	buf	destination for data
	qty	number of bytes to transfer
RETURNS	-1	on error
	numread	on success

close		
PURPOSE	Closes a file	
INCLUDE	#include <unistd.h>	
USAGE	int result = close(int fd)	
ARGS	fd	file descriptor
RETURNS	-1	on error
	0	on success

Writing who1.c

```
/* who1.c - a first version of the who program
 *          open, read UTMP file, and show results
 */
#include <stdio.h>
#include <utmp.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdlib.h>
#define SHOWHOST /* include remote machine on output */
void show_info( struct utmp* );
int main()
{
    struct utmp    current_record; /* read info into here */
    int            utmpfd;          /* read from this descriptor */
    int            reclen = sizeof(current_record);

    if ( (utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ){
        perror( UTMP_FILE ); /* UTMP_FILE is in utmp.h */
        exit(1);
    }
    while ( read(utmpfd, &current_record, reclen) == reclen )
        show_info(&current_record);
    close(utmpfd);
    return 0; /* went ok */
}
```

```

/*
 * show_info()
 *     displays contents of the utmp struct in human readable form
 *     *note* these sizes should not be hardwired
 */
void show_info( struct utmp* utbufp);
{
    printf("%-8.8s", utbufp->ut_name);    /* the logname */
    printf(" ");                          /* a space */
    printf("%-8.8s", utbufp->ut_line);    /* the tty */
    printf(" ");                          /* a space */
    printf("%10ld", utbufp->ut_time);      /* login time */
    printf(" ");                          /* a space */
#ifdef SHOWHOST
    printf("( %s)", utbufp->ut_host);      /* the host */
#endif
    printf("\n");                        /* newline */
}

```

◆ Compile and run it:

```
$ cc who1.c -o who1
```

```
$ who1
```

```
$ ./who1
```

```
system b 952601411 ()
run-level 952601411 ()
          952601416 ()
          952601416 ()
          952601417 ()
          952601417 ()
          952601419 ()
          952601419 ()
          952601423 ()
          952601566 ()
LOGIN console 952601566 ()
          tty p1 958240622 ()
shpyrko tty p2 964318862 (nas1-093.gas.swamp.org)
acotton tty p3 964319088 (math-guest04.williams.edu)
          tty p4 964320298 ()
spradlin tty p5 963881486 (h002078c6adfb.ne.rusty.net)
dkoh tty p6 964314388 (128.103.223.110)
spradlin tty p7 964058662 (h002078c6adfb.ne.rusty.net)
king tty p8 964279969 (blade-runner.mit.edu)
berschba tty p9 964188340 (dudley.learned.edu)
rserved tty p a 963538145 (gigue.eas.ivy.edu)
dabel tty p b 964319455 (roam193-27.student.state.edu)
          tty p c 964319645 ()
```

```
2007200_40634@Ubuntu-023:~/다운로드$ ./a.out
reboot ~ 1725334869 (5.15.0-94-generic)
runlevel ~ 1725334877 (5.15.0-94-generic)
2007200 127.0.0. 1725511855 (127.0.0.1:100)
```

```
2007200_40634@Ubuntu-023:~/다운로드$ ./a.out
reboot ~ 1725334869 (5.15.0-94-generic)
runlevel ~ 1725334877 (5.15.0-94-generic)
2007200 127.0.0. 1725511855 (127.0.0.1:100)
```

```
2007200_40634@Ubuntu-023:~/다운로드$ who
2007200_40634 127.0.0.1:100 2024-09-05 13:50 (127.0.0.1:100)
```

◆ Comparison

```
$ who
shpyrko ttyp2 Jul 22 22:21 (nas1-093.gas.swamp.edu)
acotton ttyp3 Jul 22 22:24 (math-guest04.williams.edu)
spradlin ttyp5 Jul 17 20:51 (h002078c6adfb.ne.rusty.net)
dkoh ttyp6 Jul 22 21:06 (128.103.223.110)
spradlin ttyp7 Jul 19 22:04 (h002078c6adfb.ne.rusty.net)
king ttyp8 Jul 22 11:32 (blade-runner.mit.edu)
berschba ttyp9 Jul 21 10:05 (dudley.learned.edu)
rserved ttypa Jul 13 21:29 (gigue.eas.ivy.edu)
dabel ttypb Jul 22 22:30 (roam193-27.student.state.edu)
rserved ttypd Jul 13 21:31 (gigue.eas.harvard.edu)
dkoh ttype Jul 22 16:46 (128.103.223.110)
molay ttyq0 Jul 22 20:03 (xyz73-200.harvard.edu)
cweiner ttyq8 Jul 21 16:40 (roam175-157.student.stats.edu)
$
```

◆ What We Need to Do

- Suppress blank records
- Get the log-in times correct

Writing who2.c

◆ Suppressing blank records

- /usr/include/utmp.h

```
/*      Definitions for ut_type      */

#define EMPTY      0
#define RUN_LVL    1
#define BOOT_TIME  2
#define OLD_TIME   3
#define NEW_TIME   4
#define INIT_PROCESS 5      /* Process spawned by "init" */
#define LOGIN_PROCESS 6     /* A "getty" process waiting for login */
#define USER_PROCESS 7     /* A user process */
#define DEAD_PROCESS 8      ※ represents the user logged into the system.
```

◆ Modification :

```
show_info( struct utmp *utbufp )
{
    if ( utbufp->ut_type != USER_PROCESS )      /* users only ! */
        return;
    printf("%-8.8s", utbufp->ut_name);           /* the username */
}
```

◆ Displaying Log-in Time in Human-Readable Form

```
$ man -k time
```

```
$ man -k time | grep transform
```

```
$ man -k time | grep -i convert
```

the **-i** option :
a *case-insensitive* match

- *How unix stores times:* **time_t**
- *Making a **time_t** readable:* **ctime**

```
$ man 3 ctime
```

```
CTIME(3)
```

```
Linux Programmer's Manual
```

```
CTIME(3)
```

```
NAME
```

```
asctime, ctime, gmtime, localtime, mktime - transform  
binary date and time to ASCII
```

```
SYNOPSIS
```

```
#include <time.h>
```

```
char *asctime(const struct tm *timeptr);
```

```
char *ctime(const time_t *timep);
```

```
...
```

```
The ctime() function converts the calendar time timep into  
a string of the form
```

```
"Wed Jun 30 21:49:08 1993\n"
```


◆ Putting it All together

```
/* who2.c - read /var/adm/utmp and list info therein
 *          - suppresses empty records
 *          - formats time nicely
 */
#include      <stdio.h>
#include      <unistd.h>
#include      <utmp.h>
#include      <fcntl.h>
#include      <time.h>
#include      <stdlib.h>
/* #define    SHOWHOST */

void showtime(long);
void show_info(struct utmp *);

int main()
{
    struct utmp    utbuf;          /* read info into here */
    int            utmpfd;         /* read from this descriptor */
```

```

    if ( (utmpfd = open(UTMP_FILE, O_RDONLY)) == -1 ){
        perror(UTMP_FILE);
        exit(1);
    }

    while( read(utmpfd, &utbuf, sizeof(utbuf)) == sizeof(utbuf) )
        show_info( &utbuf );
    close(utmpfd);
    return 0;
}

/*
 *   show info()
 *
 *       displays the contents of the utmp struct
 *       in human readable form
 *
 *       * displays nothing if record has no user name
 */
void show_info( struct utmp *utbufp )
{
    if ( utbufp->ut_type != USER_PROCESS )
        return;

    printf("%-8.8s", utbufp->ut_name);      /* the logname */
    printf(" ");                          /* a space */
    printf("%-8.8s", utbufp->ut_line);     /* the tty */
    printf(" ");                          /* a space */
    showtime( utbufp->ut_time );          /* display time */
#ifdef SHOWHOST
    if ( utbufp->ut_host[0] != '\0' )
        printf(" (%s)", utbufp->ut_host); /* the host */
#endif
    printf("\n");                        /* newline */
}

```

```

void showtime( long timeval )
/*
 *   displays time in a format fit for human consumption
 *   uses ctime to build a string then picks parts out of it
 *   Note: %12.12s prints a string 12 chars wide and LIMITS
 *   it to 12chars.
 */
{
    char    *cp;                /* to hold address of time */
    cp = ctime(&timeval);       /* convert time to string */
                                /* string looks like */
                                /* Mon Feb  4 00:46:40 EST 1991 */
                                /* 0123456789012345. */
    printf("%12.12s", cp+4 );    /* pick 12 chars from pos 4 */
}

```

Wed Jun 30 21:49:08 1993\n

partial string!

Testing who2.c

```
$ cc who2.c -o who2
```

```
$ who2
```

```
$ ./who2
```

```
rlscott  tty2    Jul 23 01:07
acotton  tty3    Jul 22 22:24
spradlin tty5    Jul 17 20:51
spradlin tty7    Jul 19 22:04
king     tty8    Jul 22 11:32
berschba tty9    Jul 21 10:05
rserverd ttya    Jul 13 21:29
rserverd ttyd    Jul 13 21:31
molay    ttyq0   Jul 22 20:03
cweiner  ttyq8    Jul 21 16:40
mnabavi  ttyx2    Apr 10 23:11
```

```
$ who
```

```
rlscott  tty2    Jul 23 01:07
acotton  tty3    Jul 22 22:24
spradlin tty5    Jul 17 20:51
spradlin tty7    Jul 19 22:04
king     tty8    Jul 22 11:32
berschba tty9    Jul 21 10:05
rserverd ttya    Jul 13 21:29
rserverd ttyd    Jul 13 21:31
molay    ttyq0   Jul 22 20:03
cweiner  ttyq8    Jul 21 16:40
mnabavi  ttyx2    Apr 10 23:11
```

```
$
```

```
2007200_40634@Ubuntu-023:~/다운로드$ ./a.out
reboot    ~          1725334869 (5.15.0-94-generic)
runlevel  ~          1725334877 (5.15.0-94-generic)
2007200_ 127.0.0. 1725511855 (127.0.0.1:100)
```

```
2007200_40634@Ubuntu-023:~/다운로드$ ./a.out
2007200_ 127.0.0. Sep  5 13:50 (127.0.0.1:100)
```

Users, Files, and the Manual

2.2 Asking about who

2.3 What Does who Do?

2.4 How Does who Do It?

2.5 Can I Write who?

2.6 Writing cp (read and write)

2.7 More Efficient File I/O: Buffering

2.8 Buffering and the Kernel

2.10 What to Do with System-Call Errors

Q1: What does `cp` do?

- ◆ `cp` makes a copy of a file

```
$ cp source-file target-file
```

- If there is no target file, `cp` *creates* it.
- If there is a target file, `cp` *replaces* the contents of that file with the contents of the source file.

Q2: How Does `cp` Create and Write?

♦ Creating/Truncating a File:

```
fd = creat(name, 0644);
```

```
open("file.txt", O_CREAT | O_WRONLY | O_TRUNC, 0644);
```

creat		
PURPOSE	Create or zero a file	
INCLUDE	#include <fcntl.h>	
USAGE	int fd = creat(char *filename, mode_t mode)	
ARGS	filename:	the name of the file
	mode:	access permission
RETURNS	-1	on error
	fd	on success

♦ Writing to a File

```
n = write(fd, buffer, num);
```

write		
PURPOSE	Send data from memory to a file	
INCLUDE	#include <unistd.h>	
USAGE	ssize_t result = write(int fd, void *buf, size_t amt)	
ARGS	fd	a file descriptor
	buf	an array
	amt	how many bytes to write
RETURNS	-1	on error
	num written	on success

Q3: Can I Write `cp`?

♦ Program outline

```
open sourcefile for reading
open copyfile   for writing
+--> read from source to buffer -- eof? --+
|___ write from buffer to copy             |
                                          |
close sourcefile      <-----+
close copyfile
```




FIGURE 2.4

Copying a file by reading and writing.

```

/** cpl.c
 *      version 1 of cp - uses read and write with tunable buffer size
 *
 *      usage: cpl src dest
 */
#include      <stdio.h>
#include      <unistd.h>
#include      <fcntl.h>
#include      <stdlib.h>
#define BUFFERSIZE      4096
#define COPYMODE      0644

void oops(char *, char *);

main(int ac, char *av[])
{
    int      in_fd, out_fd, n_chars;
    char      buf[BUFFERSIZE];

                                /* check args */
    if ( ac != 3 ){
        fprintf( stderr, "usage: %s source destination\n", *av);
        exit(1);
    }

                                /* open files */
    if ( (in_fd=open(av[1], O_RDONLY)) == -1 )
        oops("Cannot open ", av[1]);

    if ( (out_fd=creat( av[2], COPYMODE)) == -1 )
        oops( "Cannot creat", av[2]);

                                /* copy files */

```

```

while ( (n_chars = read(in_fd , buf, BUFFERSIZE)) > 0 )
    if ( write( out_fd, buf, n_chars ) != n_chars )
        oops("Write error to ", av[2]);
if ( n_chars == -1 )
    oops("Read error from ", av[1]);

                                /* close files */

if ( close(in_fd) == -1 || close(out_fd) == -1 )
    oops("Error closing files","");
}

void oops(char *s1, char *s2)
{
    fprintf(stderr, "Error: %s ", s1);
    perror(s2);
    exit(1);
}

```

```
$ cc cp1.c -o cp1
$ cp1 cp1 copy.of.cp1    ✂ $ ./cp1 cp1 copy.of.cp1
$ ls -l cp1 copy.of.cp1
-rw-r--r--    1 bruce    bruce          37419 Jul 23 03:12 copy.of.cp1
-rwxrwxr-x    1 bruce    bruce          37419 Jul 23 03:08 cp1
$ cmp cp1 copy.of.cp1    ✂ $ man cmp
$
```

Users, Files, and the Manual

2.2 Asking about who

2.3 What Does who Do?

2.4 How Does who Do It?

2.5 Can I Write who?

2.6 Writing cp (read and write)

2.7 More Efficient File I/O: Buffering

2.8 Buffering and the Kernel


2.10 What to Do with System-Call Errors

Does the Size of the Buffer Matter?

♦ Execution time for cp1

- `read(in_fd, buf, BUFFERSIZE);`

buffersize	execution time in seconds
1	50.29
4	12.81
16	3.28
64	0.96
128	0.56
256	0.37
512	0.27
1024	0.22
2048	0.19
4096	0.18
8192	0.18
16384	0.18



Ex: Filesize = 2500 bytes

If buffer = 100 bytes then
copy requires 25 read() and 25 write() calls

If buffer = 1000 bytes then
copy requires 3 read() and 3 write() calls

Why System Calls are Time Consuming?

- ◆ **Not only does transferring data take time but **Mode Change** takes time**
 - It runs various Kernel functions
 - It also requires a shift from USER MODE to KERNEL MODE and back;
- ◆ Thus, try to **minimize system calls**




FIGURE 2.5
Control flow during system calls.

who2.c is Inefficient?

- ♦ **who2.c** use one **system call** for each **utmp** record

```
struct utmp utbuf;           /* read info into here */
int utmpfd;                  /* read from this descriptor */

if( ( utmpfd = open(UTMP_FILE, O_RDONLY) ) == -1 ) {
    perror( UTMP_FILE );     /* UTMP_FILE is in utmp.h */
    exit(1);
}

while( read(utmpfd, &utbuf, sizeof(utbuf) ) == sizeof(utbuf) )
    show_info(&utbuf);
close(utmpfd);
```

- A better idea: ...

Adding Buffering : who3.c

File buffering with utmplib




FIGURE 2.6

Buffering disk data in user space.

main calls a function in `utmplib.c` to get the next struct `utmp`.

Functions in `utmplib.c` read structs 16 at a time from the disk into an array.

The kernel is called only when all 16 are used up.

```
$ gcc who3.c utmplib.c -o who3
$ ./who3
```

```
/* who3.c - who with buffered reads
 *         - surpresses empty records
 *         - formats time nicely
 *         - buffers input (using utmp/lib)
 */
```

```
#include <stdio.h>
#include <sys/types.h>
#include <utmp.h>
#include <fcntl.h>
#include <time.h>
#include <stdlib.h>
#define SHOWHOST
```

```
void show_info(struct utmp *);
void showtime(time_t);
```

```
int main()
{
```

```
    struct utmp    *utbufp,          /* holds pointer to next rec    */
                  *utmp_next();      /* returns pointer to next     */
```

```
    if ( utmp_open( UTMP_FILE ) == -1 ){
        perror(UTMP_FILE);
        exit(1);
    }
```

```
    while ( ( utbufp = utmp_next() ) != ((struct utmp *) NULL) )
        show_info( utbufp );
    utmp_close( );
    return 0;
```

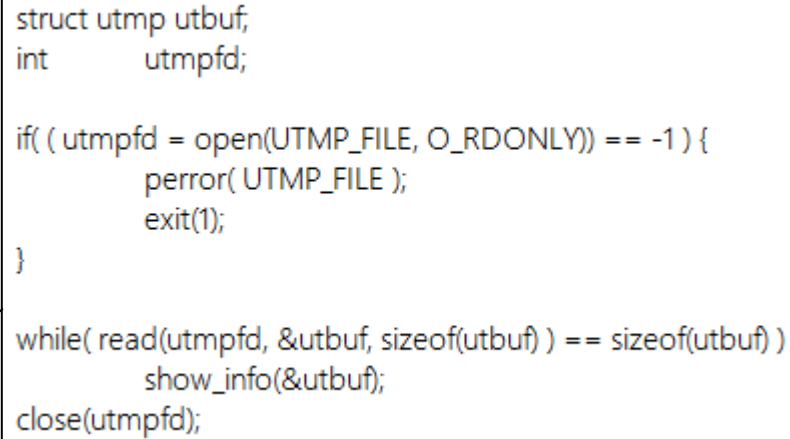
```
}
```

```
/*
 * show info()
 * ...
```

```
struct utmp utbuf;
int utmpfd;

if( ( utmpfd = open(UTMP_FILE, O_RDONLY) ) == -1 ) {
    perror( UTMP_FILE );
    exit(1);
}

while( read(utmpfd, &utbuf, sizeof(utbuf) ) == sizeof(utbuf) )
    show_info(&utbuf);
close(utmpfd);
```



```
/* utmplib.c - functions to buffer reads from utmp file
*
*     functions are
*         utmp_open( filename )    - open file
*                                     returns -1 on error
*         utmp_next( )             - return pointer to next struct
*                                     returns NULL on eof
*         utmp_close()             - close file
*
*     reads NRECS per read and then doles them out from the buffer
*/
```

```

#include      <stdio.h>
#include      <fcntl.h>
#include      <sys/types.h>
#include      <utmp.h>

#define NRECS 16
#define NULLUT ((struct utmp *)NULL)
#define UTSIZE (sizeof(struct utmp))

static char utmpbuf[NRECS * UTSIZE];          /* storage */
static int  num_recs;                         /* num stored */
static int  cur_rec;                          /* next to go */
static int  fd_utmp = -1;                    /* read from */

utmp_open( char *filename )
{
    fd_utmp = open( filename, O_RDONLY );      /* open it */
    cur_rec = num_recs = 0;                  /* no recs yet */
    return fd_utmp;                          /* report */
}

struct utmp *utmp_next()
{
    struct utmp *recp;
    if ( fd_utmp == -1 )                      /* error ? */
        return NULLUT;
    if ( cur_rec==num_recs && utmp_reload()==0 ) /* any more ? */
        return NULLUT;
                                           /* get address of next record */
    recp = ( struct utmp *) &utmpbuf[cur_rec * UTSIZE];
    cur_rec++;
    return recp;
}

```

```

int utmp_reload()
/*
 *   read next bunch of records into buffer
 */
{
    int      amt_read;

                                /* read them in          */
    amt_read = read( fd_utm , utmpbuf, NRECS * UTSIZE );

                                /* how many did we get? */
    num_recs = amt_read/UTSIZE;

                                /* reset pointer          */
    cur_rec  = 0;
    return num_recs;
}

```

```

utmp_close()
{
    if ( fd_utm != -1 )          /* don't close if not */
        close( fd_utm );        /* open                */
}

```

Users, Files, and the Manual

2.2 Asking about who

2.3 What Does who Do?

2.4 How Does who Do It?

2.5 Can I Write who?

2.6 Writing cp (read and write)

2.7 More Efficient File I/O: Buffering

2.8 Buffering and the Kernel

2.10 What to Do with System-Call Errors

If Buffering Is So Smart, Why Doesn't the Kernel Do It?

◆ It does!

- **Kernel buffer** keeps copies of disk block in memory
- The `read()` call copies data into a process **from a kernel buffer**
- The `write()` copies data from the process **to a kernel buffer**
- if not in a kernel buffer? ...




FIGURE 2.7
Buffering disk data in the kernel.

Consequences of Kernel Buffering

- Faster “disk” I/O
- Optimized disk writes
- Need to write buffers to disk before shutdown




FIGURE 2.7

Buffering disk data in the kernel.

Users, Files, and the Manual

2.2 Asking about who

2.3 What Does who Do?

2.4 How Does who Do It?

2.5 Can I Write who?

2.6 Writing cp (read and write)

2.7 More Efficient File I/O: Buffering

2.8 Buffering and the Kernel

2.10 What to Do with System-Call Errors

What to Do with System-Call Errors

◆ System Calls and Errors:

- open, read, and lseek return -1 when an error occurs.

◆ How to identify what went wrong : **errno**

- The kernel tells your program the cause of the error *by storing an global variable called **errno***.
- Every program contains this variable.

◆ Common error codes include:

\$ man 3 errno

```
#define EPERM          1      /* Operation not permitted */
#define ENOENT          2      /* No such file or directory */
#define ESRCH          3      /* No such process */
#define EINTR          4      /* Interrupted system call */
#define EIO            5      /* I/O error */
```

```
#include <errno.h>

extern int errno;

int sample()
{
    int fd;
    fd = open("file", O_RDONLY);
    if ( fd == -1 )
    {
        printf("Cannot open file: ");
        if ( errno == ENOENT )
            printf("There is no such file.");
        else if ( errno == EINTR )
            printf("Interrupted while opening file.");

        else if ( errno == EACCESS )
            printf("You do not have permission to open file.");
        ...
    }
}
```

```
1 #include <stdio.h>
2 #include <errno.h>
3 #include <string.h>
4 #include <fcntl.h>
5
6 int main() {
7     int fd;
8
9     fd = open("nonexistentfile.txt", O_RDONLY);
10
11     if (fd == -1) {
12         printf("Error %d: %s\n", errno, strerror(errno));
13     } else {
14         printf("File opened.\n");
15     }
16
17     return 0;
18 }
```

```
2007200_40634@Ubuntu-037:~/다운로드$ ./a.out
Error 2: No such file or directory
2007200_40634@Ubuntu-037:~/다운로드$
```

◆ Reporting Errors: \$ man 3 perror

- Print a system error message

```
void perror(const char *s);  
char *strerror(int errnum);
```

```
int sample()  
{  
    int fd;  
    fd = open("file", O_RDONLY);  
    if ( fd == -1 )  
    {  
        perror("Cannot open file");  
        return;  
    }  
    ...  
}
```

Users, Files, and the Manual

2.2 Asking about **who**

2.3 What Does **who** Do?

2.4 How Does **who** Do It?

2.5 Can I Write **who**?

2.6 Writing **cp** (read and write)

2.7 More Efficient File I/O: **Buffering**

2.8 Buffering and the Kernel

2.10 What to Do with **System-Call Errors**